

Logistic Regression project

December 1, 2024

1 TikTok Project

2 Objectives

1. **Develop a Machine Learning Model:** Create a predictive model to classify TikTok videos as containing either a claim or an opinion.
2. **Enhance Efficiency in Handling User Reports:** Use the predictive model to reduce the backlog of user reports by prioritizing them based on content classification.
3. **Complete Project Development Steps:** Finalize ongoing work, including:
 - Initial action plan creation.
 - Python-based initial coding.
 - Exploratory Data Analysis (EDA).
 - Hypothesis testing.
4. **Analyze Hypothesis Testing Results:** Evaluate the relationship between user variables and verified status based on hypothesis testing.
5. **Explore Verified User Patterns:** Investigate how verified users are associated with posting opinions to refine model predictions.
6. **Conduct Logistic Regression Analysis:** Use verified status as the outcome variable to understand video characteristics linked to user verification.
7. **Refine the Final Model:** Incorporate logistic regression insights into the final machine learning model to improve claim vs opinion predictions.
8. **Support TikTok's Operations Team:** Provide actionable insights a Abadi and the operations team to enhance understanding of video characteristics and user behavior patterns.

3 Steps

1. **Demonstrate Knowledge of EDA and Regression Models:** Showcase your understanding of Exploratory Data Analysis (EDA) and logistic regression by building and evaluating a predictive model.

2. **Build a Logistic Regression Model:** Develop a logistic regression model to estimate the probability of a specific outcome using Python.
3. **Evaluate the Logistic Regression Model:** Assess the model's performance and ensure it meets the assumptions required for logistic regression-
4. **Part 1: EDA & Checking Model Assumptions:**
 - Understand the purpose of EDA before constructing a logistic regression model.
 - Identify patterns, relationships, and potential issues in the data-t.
5. **Part 2: Model Building and Evaluation:**
 - Utilize appropriate resources and methodologies to construct and evaluate the logistic regression model.
 - Document challenges and resources used during the p-cess.
6. **Part 3: Interpreting Model Results:**
 - Extract key insights from the logistic regression model.
 - Formulate business recommendations based on model results.
7. **Write an Executive Summary:**
 - Use the PACE Strategy Document to summarize findings, insights, and business recommendations effectively.
8. **Prepare for Comparison:** Complete the activity to compare your work with a provided exemplar for further learning and refinement.

3.0.1 Task 1. Imports and loading

Import the data and packages for building Logistic regression models.

```
[9]: # Import packages for data manipulation
import pandas as pd
import numpy as np
# Import packages for data visualization
import seaborn as sns
import matplotlib.pyplot as plt
# Import packages for data preprocessing
from sklearn.preprocessing import OneHotEncoder
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
from sklearn.utils import resample
# Import packages for data modeling
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
```

Load the TikTok dataset.

```
[10]: # Load dataset into dataframe
data = pd.read_csv("tiktok_dataset.csv")
```

EDA is the important step in data analytics, we can see the descriptive statistic, outlier, dataset info, data with duplicate and data with missing values. In Logistic Regression, we need to see the correlation between variable, to verify model assumptions such as no severe multicollinearity.

3.0.2 Task 2a. Explore data with EDA

Analyze the data and check for and handle missing values and duplicates.

```
[11]: # Display first few rows
data.head()
```

```
[11]:  # claim_status    video_id  video_duration_sec  \
0  1          claim  7017666017             59
1  2          claim  4014381136             32
2  3          claim  9859838091             31
3  4          claim  1866847991             25
4  5          claim  7105231098             19

                                video_transcription_text  verified_status  \
0  someone shared with me that drone deliveries a...    not verified
1  someone shared with me that there are more mic...    not verified
2  someone shared with me that american industria...    not verified
3  someone shared with me that the metro of st. p...    not verified
4  someone shared with me that the number of busi...    not verified

author_ban_status  video_view_count  video_like_count  video_share_count  \
0      under review         343296.0         19425.0           241.0
1           active         140877.0         77355.0          19034.0
2           active         902185.0         97690.0           2858.0
3           active         437506.0        239954.0          34812.0
4           active         56167.0         34987.0           4110.0

video_download_count  video_comment_count
0              1.0              0.0
1            1161.0             684.0
2              833.0             329.0
3            1234.0             584.0
4              547.0             152.0
```

Get the number of rows and columns in the dataset.

```
[12]: # Get number of rows and columns
print(f"Total rows : {data.shape[0]}")
print(f"Total columns : {data.shape[1]}")
print(f"Size : {data.size}")
```

Total rows : 19382
Total columns : 12
Size : 232584

Get the data types of the columns.

```
[13]: # Get data types
data.dtypes
```

```
[13]: #                int64
claim_status      object
video_id          int64
video_duration_sec int64
video_transcription_text object
verified_status   object
author_ban_status object
video_view_count  float64
video_like_count  float64
video_share_count float64
video_download_count float64
video_comment_count float64
dtype: object
```

Get basic information about the dataset.

```
[14]: # Basic information
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19382 entries, 0 to 19381
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   #                      19382 non-null  int64
1   claim_status           19084 non-null  object
2   video_id               19382 non-null  int64
3   video_duration_sec     19382 non-null  int64
4   video_transcription_text 19084 non-null  object
5   verified_status        19382 non-null  object
6   author_ban_status      19382 non-null  object
7   video_view_count       19084 non-null  float64
8   video_like_count       19084 non-null  float64
9   video_share_count      19084 non-null  float64
10  video_download_count    19084 non-null  float64
11  video_comment_count     19084 non-null  float64
dtypes: float64(5), int64(3), object(4)
memory usage: 1.8+ MB
```

Generate basic descriptive statistics about the dataset.

```
[15]: # Generate basic descriptive stats
data.describe()
```

```
[15]:
```

	#	video_id	video_duration_sec	video_view_count	\
count	19382.000000	1.938200e+04	19382.000000	19084.000000	
mean	9691.500000	5.627454e+09	32.421732	254708.558688	
std	5595.245794	2.536440e+09	16.229967	322893.280814	
min	1.000000	1.234959e+09	5.000000	20.000000	
25%	4846.250000	3.430417e+09	18.000000	4942.500000	
50%	9691.500000	5.618664e+09	32.000000	9954.500000	
75%	14536.750000	7.843960e+09	47.000000	504327.000000	
max	19382.000000	9.999873e+09	60.000000	999817.000000	

	video_like_count	video_share_count	video_download_count	\
count	19084.000000	19084.000000	19084.000000	
mean	84304.636030	16735.248323	1049.429627	
std	133420.546814	32036.174350	2004.299894	
min	0.000000	0.000000	0.000000	
25%	810.750000	115.000000	7.000000	
50%	3403.500000	717.000000	46.000000	
75%	125020.000000	18222.000000	1156.250000	
max	657830.000000	256130.000000	14994.000000	

	video_comment_count
count	19084.000000
mean	349.312146
std	799.638865
min	0.000000
25%	1.000000
50%	9.000000
75%	292.000000
max	9599.000000

```
[16]: # Object Variable
data.describe(include='object').T
```

```
[16]:
```

	count	unique	\
claim_status	19084	2	
video_transcription_text	19084	19012	
verified_status	19382	2	
author_ban_status	19382	3	

	top	\
claim_status	claim	
video_transcription_text	a friend read in the media a claim that badmi...	
verified_status	not verified	
author_ban_status	active	

	freq
claim_status	9608
video_transcription_text	2
verified_status	18142
author_ban_status	15663

Check for and handle missing values.

```
[17]: # Check for missing values
data.isna().sum()
```

```
[17]: #
claim_status      298
video_id          0
video_duration_sec 0
video_transcription_text 298
verified_status    0
author_ban_status  0
video_view_count   298
video_like_count   298
video_share_count  298
video_download_count 298
video_comment_count 298
dtype: int64
```

```
[18]: data.loc[data.isna().any(axis=1)].head()
```

```
[18]:      # claim_status  video_id  video_duration_sec  \
19084  19085      NaN  4380513697      39
19085  19086      NaN  8352130892      60
19086  19087      NaN  4443076562      25
19087  19088      NaN  8328300333       7
19088  19089      NaN  3968729520       8

      video_transcription_text  verified_status  author_ban_status  \
19084      NaN      not verified      active
19085      NaN      not verified      active
19086      NaN      not verified      active
19087      NaN      not verified      active
19088      NaN      not verified      active

      video_view_count  video_like_count  video_share_count  \
19084      NaN      NaN      NaN
19085      NaN      NaN      NaN
19086      NaN      NaN      NaN
19087      NaN      NaN      NaN
```

19088	NaN	NaN	NaN
-------	-----	-----	-----

	video_download_count	video_comment_count
19084	NaN	NaN
19085	NaN	NaN
19086	NaN	NaN
19087	NaN	NaN
19088	NaN	NaN

```
[19]: # Drop rows with missing values
data.dropna(axis=0,inplace=True)
```

```
[20]: # Display first few rows after handling missing values
data.head()
```

```
[20]: # claim_status    video_id  video_duration_sec  \
0  1      claim    7017666017          59
1  2      claim    4014381136          32
2  3      claim    9859838091          31
3  4      claim    1866847991          25
4  5      claim    7105231098          19

      video_transcription_text  verified_status  \
0  someone shared with me that drone deliveries a...  not verified
1  someone shared with me that there are more mic...  not verified
2  someone shared with me that american industria...  not verified
3  someone shared with me that the metro of st. p...  not verified
4  someone shared with me that the number of busi...  not verified

      author_ban_status  video_view_count  video_like_count  video_share_count  \
0      under review      343296.0      19425.0      241.0
1      active      140877.0      77355.0      19034.0
2      active      902185.0      97690.0      2858.0
3      active      437506.0      239954.0      34812.0
4      active      56167.0      34987.0      4110.0

      video_download_count  video_comment_count
0          1.0          0.0
1        1161.0         684.0
2          833.0         329.0
3        1234.0         584.0
4          547.0         152.0
```

Check for and handle duplicates.

```
[21]: # Check for duplicates
data.duplicated().sum()
```

[21]: 0


Check for and handle outliers.

```
[22]: # Create a boxplot to visualize distribution of
#
# ↳ `video_duration_sec`, `video_view_count`, `video_like_count`, `video_comment_count`

fig, ax = plt.subplots(1,4, figsize=(12,1.5), constrained_layout=True)
fig.suptitle('Boxplots of Distribution')

sns.boxplot(data=data, x='video_duration_sec', ax=ax[0], width=.5, fliersize=3)
sns.boxplot(data=data, x='video_view_count', ax=ax[1], width=.5, fliersize=3)
sns.boxplot(data=data, x='video_like_count', ax=ax[2], width=.5, fliersize=3)
sns.boxplot(data=data, x='video_comment_count', ax=ax[3], width=.5, fliersize=3)

plt.show()
```



```
[26]: # Outliers imputation with IQR
def imputation(df, columns, iqr_factor):
    for col in columns:
        # Q1 - Q3
        q1 = df[col].quantile(0.25)
        q3 = df[col].quantile(0.75)
        # IQR
        iqr = q3 - q1
        # upper limit
        upper_limit = q3 + (iqr_factor * iqr)
        # reassign
        df.loc[df[col] > upper_limit, col] = upper_limit

    print(f"Columns : {col}")
    print(f"Upper Limit : {upper_limit}\n")
    print(f"After : ")
    print(df[col].describe())
```

```
[27]: # Check for and handle outliers for video_like_count
data['video_like_count'].describe()
```



```
[27]: count      19084.000000
      mean       84304.636030
      std        133420.546814
      min         0.000000
      25%         810.750000
      50%         3403.500000
      75%        125020.000000
      max        657830.000000
      Name: video_like_count, dtype: float64
```

```
[28]: # video_like_count imputation
      imputation(data,['video_like_count'],1.5)
```

```
Columns : video_like_count
Upper Limit : 311333.875
```

```
After :
count      19084.000000
mean       74323.538632
std        107103.555220
min         0.000000
25%         810.750000
50%         3403.500000
75%        125020.000000
max        311333.875000
Name: video_like_count, dtype: float64
```

```
[29]: # video_comment_count
      data['video_comment_count'].describe()
```

```
[29]: count      19084.000000
      mean       349.312146
      std        799.638865
      min         0.000000
      25%         1.000000
      50%         9.000000
      75%        292.000000
      max        9599.000000
      Name: video_comment_count, dtype: float64
```

```
[30]: # video_comment_count imputation
      imputation(data,['video_comment_count'],1.5)
```

```
Columns : video_comment_count
Upper Limit : 728.5
```

```
After :
count      19084.000000
```

```

mean      181.023501
std       272.084766
min        0.000000
25%        1.000000
50%        9.000000
75%       292.000000
max       728.500000
Name: video_comment_count, dtype: float64

```

Check class balance of the target variable. Remember, the goal is to predict whether the user of a given post is verified or unverified.

```

[31]: # Check class balance
data['verified_status'].value_counts(normalize=True)*100

```

```

[31]: verified_status
not verified    93.71201
verified        6.28799
Name: proportion, dtype: float64

```

Approximately 93.7% of the dataset represents videos posted by unverified accounts and 6.3% represents videos posted by verified accounts. So the outcome variable is not very balanced.

Use resampling to create class balance in the outcome variable, if needed.

```

[32]: # Use resampling to create class balance in the outcome variable

# Identify data points from majority and minority classes
data_majority = data.loc[data['verified_status'] == 'not verified']
data_minority = data.loc[data['verified_status'] == 'verified']

# Upsample the minority class ("verified")
minority_upsampled =
↳resample(data_minority,n_samples=len(data_majority),replace=True,↳
↳random_state=0)

# Combine majority class with upsampled minority class
data_upsampled = pd.concat([data_majority,minority_upsampled]).
↳reset_index(drop=True)

# Display new class counts
data_upsampled['verified_status'].value_counts()

```

```

[32]: verified_status
not verified    17884
verified        17884
Name: count, dtype: int64

```

```
[33]: # new class percentage
print('Percentage :')
data_upsampled['verified_status'].value_counts(normalize=True)* 100
```

Percentage :

```
[33]: verified_status
not verified    50.0
verified       50.0
Name: proportion, dtype: float64
```

Get the average video_transcription_text length for videos posted by verified accounts and the average video_transcription_text length for videos posted by unverified accounts.

```
[34]: # Get the average `video_transcription_text` length for claims and the average
      ↳ `video_transcription_text` length for opinions
verified_and_video_transcription_text =
      ↳ data_upsampled[['verified_status', 'video_transcription_text']]
verified_and_video_transcription_text.
      ↳ groupby(['verified_status'])[['video_transcription_text']].agg(func= lambda
      ↳ x : np.mean([len(text) for text in x ]))
```

```
[34]:          video_transcription_text
verified_status
not verified    89.401141
verified       84.569559
```

Extract the length of each video_transcription_text and add this as a column to the dataframe, so that it can be used as a potential feature in the model.

```
[35]: # Extract the length of each `video_transcription_text` and add this as a
      ↳ column to the dataframe
data_upsampled['text_length'] = data_upsampled['video_transcription_text'].
      ↳ apply(func=lambda x : len(x))
```

```
[36]: # Display first few rows of dataframe
data_upsampled.head()
```

```
[36]:  # claim_status    video_id  video_duration_sec  \
0  1      claim    7017666017          59
1  2      claim    4014381136          32
2  3      claim    9859838091          31
3  4      claim    1866847991          25
4  5      claim    7105231098          19

          video_transcription_text  verified_status  \
0  someone shared with me that drone deliveries a...  not verified
1  someone shared with me that there are more mic...  not verified
```

2	someone shared with me that american industria...	not verified
3	someone shared with me that the metro of st. p...	not verified
4	someone shared with me that the number of busi...	not verified

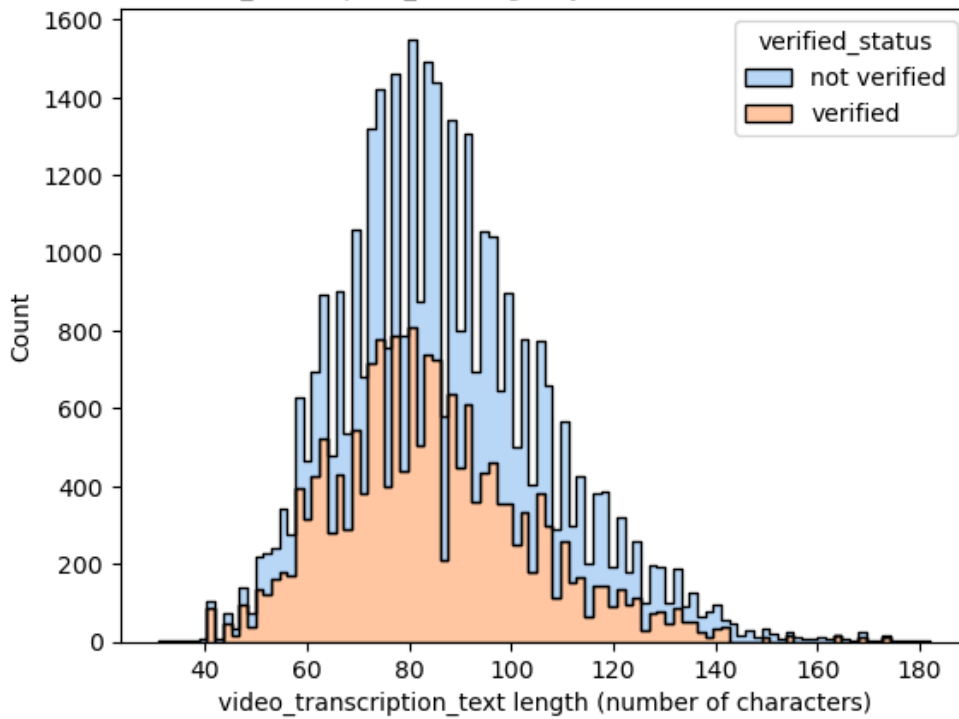
	author_ban_status	video_view_count	video_like_count	video_share_count	\
0	under review	343296.0	19425.0	241.0	
1	active	140877.0	77355.0	19034.0	
2	active	902185.0	97690.0	2858.0	
3	active	437506.0	239954.0	34812.0	
4	active	56167.0	34987.0	4110.0	

	video_download_count	video_comment_count	text_length
0	1.0	0.0	97
1	1161.0	684.0	107
2	833.0	329.0	137
3	1234.0	584.0	131
4	547.0	152.0	128

Visualize the distribution of video_transcription_text length for videos posted by verified accounts and videos posted by unverified accounts.

```
[37]: # Visualize the distribution of `video_transcription_text` length by verified
      ↪accounts and unverified accounts
      # Create two histograms in one plot
      sns.histplot(data=data_upsampled,
      ↪x='text_length',hue='verified_status',multiple='stack',element='step',palette='pastel')
      plt.xlabel('video_transcription_text length (number of characters)')
      plt.title('Distribution of video_transcription_text length by verified accounts,
      ↪and unverified accounts', fontsize=10)
      plt.show()
```

Distribution of video_transcription_text length by verified accounts and unverified accounts



3.0.3 Task 2b. Examine correlations

Next, code a correlation matrix to help determine most correlated variables.

```
[38]: # determine most correlated variables
correlation = data_upsampled.corr(numeric_only=True, method='pearson')
correlation
```

```
[38]:
```

	#	video_id	video_duration_sec	\
#	1.000000	-0.000853	-0.011729	
video_id	-0.000853	1.000000	0.011859	
video_duration_sec	-0.011729	0.011859	1.000000	
video_view_count	-0.697007	0.002554	0.013589	
video_like_count	-0.626385	0.005993	0.004494	
video_share_count	-0.504015	0.010515	0.002206	
video_download_count	-0.487096	0.008753	0.003989	
video_comment_count	-0.608773	0.012674	-0.001086	
text_length	-0.193677	-0.007083	-0.002981	

	video_view_count	video_like_count	video_share_count	\
#	-0.697007	-0.626385	-0.504015	
video_id	0.002554	0.005993	0.010515	
video_duration_sec	0.013589	0.004494	0.002206	

video_view_count	1.000000	0.856937	0.711313
video_like_count	0.856937	1.000000	0.832146
video_share_count	0.711313	0.832146	1.000000
video_download_count	0.690048	0.805543	0.710117
video_comment_count	0.748361	0.818032	0.671335
text_length	0.244693	0.216693	0.171651

	video_download_count	video_comment_count	text_length
#	-0.487096	-0.608773	-0.193677
video_id	0.008753	0.012674	-0.007083
video_duration_sec	0.003989	-0.001086	-0.002981
video_view_count	0.690048	0.748361	0.244693
video_like_count	0.805543	0.818032	0.216693
video_share_count	0.710117	0.671335	0.171651
video_download_count	1.000000	0.793668	0.173396
video_comment_count	0.793668	1.000000	0.217661
text_length	0.173396	0.217661	1.000000

```
[39]: correlation.style.background_gradient('crest')
```

```
[39]: <pandas.io.formats.style.Styler at 0x76c78dae8850>
```

Visualize a correlation heatmap of the data.

```
[40]: # Create a heatmap to visualize how correlated variables are

plt.figure(figsize=(10,5))
sns.heatmap(data_upsampled[['video_duration_sec', 'video_view_count',
↪ 'video_like_count',
    'video_share_count', 'video_download_count', 'video_comment_count',
    'text_length']].corr(),annot=True,cmap='crest')
```

```
[40]: <Axes: >
```



One of the model assumptions for logistic regression is no severe multicollinearity among the features. Take this into consideration as you examine the heatmap and choose which features to proceed with.

The above heatmap shows that the following pair of variables are strongly correlated: video_view_count and video_like_count (0.86 correlation coefficient).

3.0.4 Task 3a. Select variables

Set your Y and X variables.

```
[75]: # Select outcome variable
y = data_upsampled['verified_status']
```

Select the features.

```
[76]: # Select features
X = data_upsampled[["video_duration_sec", "claim_status", "author_ban_status",
↪ "video_view_count", "video_share_count", "video_download_count",
↪ "video_comment_count"]]

# Display first few rows of features dataframe
X.head()
```

```
[76]: video_duration_sec claim_status author_ban_status video_view_count \
0          59          claim      under review      343296.0
1          32          claim          active      140877.0
2          31          claim          active      902185.0
3          25          claim          active      437506.0
4          19          claim          active      56167.0

      video_share_count video_download_count video_comment_count
0          241.0          1.0          0.0
1       19034.0       1161.0       684.0
2        2858.0         833.0       329.0
3       34812.0       1234.0       584.0
4        4110.0         547.0       152.0
```

3.0.5 Task 3b. Train-test split

Split the data into training and testing sets.

```
[77]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.25,
↳ random_state=0)
```

Confirm that the dimensions of the training and testing sets are in alignment.

```
[80]: # Get shape of each training and testing set
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
[80]: ((26826, 7), (8942, 7), (26826,), (8942,))
```

3.0.6 Task 3c. Encode variables

Check the data types of the features.

```
[82]: # Check data types
X_train.dtypes
```

```
[82]: video_duration_sec      int64
claim_status                object
author_ban_status           object
video_view_count            float64
video_share_count           float64
video_download_count        float64
video_comment_count         float64
dtype: object
```

```
[83]: # Get unique values in `claim_status`
X_train['claim_status'].unique()
```



```
[83]: array(['opinion', 'claim'], dtype=object)
```

```
[84]: # Get unique values in `author_ban_status`  
X_train['author_ban_status'].unique()
```

```
[84]: array(['active', 'under review', 'banned'], dtype=object)
```

As shown above, the `claim_status` and `author_ban_status` features are each of data type `object` currently. In order to work with the implementations of models through `sklearn`, these categorical features will need to be made numeric. One way to do this is through one-hot encoding.

Encode categorical features in the training set using an appropriate method.

```
[85]: # Select the training features that needs to be encoded  
X_train_encode = X_train[['claim_status', 'author_ban_status']]  
  
# Display first few rows  
X_train_encode
```

```
[85]:      claim_status  author_ban_status  
33058      opinion          active  
20491      opinion          active  
25583      opinion          active  
18474      opinion          active  
27312      opinion          active  
...      ...      ...  
20757      opinion          active  
32103      opinion          active  
30403      opinion          active  
21243      opinion          active  
2732      claim          banned
```

```
[26826 rows x 2 columns]
```

```
[86]: # Set up an encoder for one-hot encoding the categorical features  
X_encoder = OneHotEncoder(drop='first', sparse_output=False)
```

```
[87]: # Fit and transform the training features using the encoder  
X_train_encoded = X_encoder.fit_transform(X_train_encode)
```

```
[89]: # Get feature names from encoder  
X_encoder.get_feature_names_out()
```

```
[89]: array(['claim_status_opinion', 'author_ban_status_banned',  
        'author_ban_status_under review'], dtype=object)
```

```
[90]: # Display first few rows of encoded training features  
X_train_encoded
```

```
[90]: array([[1., 0., 0.],
           [1., 0., 0.],
           [1., 0., 0.],
           ...,
           [1., 0., 0.],
           [1., 0., 0.],
           [0., 1., 0.]])
```

```
[91]: # Place encoded training features (which is currently an array) into a dataframe
X_train_encoded_df = pd.DataFrame(X_train_encoded, columns=X_encoder.
    ↪get_feature_names_out())

# Display first few rows
X_train_encoded_df
```

```
[91]:
```

	claim_status_opinion	author_ban_status_banned \
0	1.0	0.0
1	1.0	0.0
2	1.0	0.0
3	1.0	0.0
4	1.0	0.0
...
26821	1.0	0.0
26822	1.0	0.0
26823	1.0	0.0
26824	1.0	0.0
26825	0.0	1.0

	author_ban_status_under review
0	0.0
1	0.0
2	0.0
3	0.0
4	0.0
...	...
26821	0.0
26822	0.0
26823	0.0
26824	0.0
26825	0.0

```
[26826 rows x 3 columns]
```

```
[92]: # Display first few rows of `X_train` with `claim_status` and
    ↪`author_ban_status` columns dropped (since these features are being
    ↪transformed to numeric)
X_train.drop(columns=["claim_status", "author_ban_status"]).head()
```

```
[92]:      video_duration_sec  video_view_count  video_share_count  \
33058          33          2252.0          23.0
20491          52          6664.0          550.0
25583          37          6327.0          257.0
18474          57          1702.0          28.0
27312          21          3842.0          101.0

      video_download_count  video_comment_count
33058          4.0          0.0
20491         53.0          2.0
25583          3.0          0.0
18474          0.0          0.0
27312          1.0          0.0
```

```
[93]: # Concatenate `X_train` and `X_train_encoded_df` to form the final dataframe
      ↪for training data (`X_train_final`)
X_train_final = pd.concat([X_train.drop(columns=["claim_status",
      ↪"author_ban_status"]), X_train_encoded_df], axis=1)

# Display first few rows
X_train_final
```

```
[93]:      video_duration_sec  video_view_count  video_share_count  \
0          33          2252.0          23.0
1          52          6664.0          550.0
2          37          6327.0          257.0
3          57          1702.0          28.0
4          21          3842.0          101.0
...
26821          36          8848.0          441.0
26822          25          8821.0          134.0
26823          26          958.0          21.0
26824          32          8553.0          744.0
26825          47          484238.0          6432.0

      video_download_count  video_comment_count  claim_status_opinion  \
0          4.0          0.0          1.0
1         53.0          2.0          1.0
2          3.0          0.0          1.0
3          0.0          0.0          1.0
4          1.0          0.0          1.0
...
26821         24.0          1.0          1.0
26822          8.0          1.0          1.0
26823          2.0          0.0          1.0
26824         62.0         23.0          1.0
26825        104.0          1.0          0.0
```

	author_ban_status_banned	author_ban_status_under review
0	0.0	0.0
1	0.0	0.0
2	0.0	0.0
3	0.0	0.0
4	0.0	0.0
...
26821	0.0	0.0
26822	0.0	0.0
26823	0.0	0.0
26824	0.0	0.0
26825	1.0	0.0

[26826 rows x 8 columns]

Check the data type of the outcome variable.

```
[94]: # Check data type of outcome variable
y_train.dtypes
```

```
[94]: dtype('O')
```

```
[95]: # Get unique values of outcome variable
y_train.unique()
```

```
[95]: array(['verified', 'not verified'], dtype=object)
```

As shown above, the outcome variable is of data type `object` currently. One-hot encoding can be used to make this variable numeric.

Encode categorical values of the outcome variable the training set using an appropriate method.

```
[97]: # Set up an encoder for one-hot encoding the categorical outcome variable
y_encoder = OneHotEncoder(drop='first', sparse_output=False)
```

```
[98]: # Encode the training outcome variable
# - Adjusting the shape of `y_train` before passing into `.fit_transform()`,
#   ↳ since it takes in 2D array
# - Using `.ravel()` to flatten the array returned by `.fit_transform()`, so
#   ↳ that it can be used later to train the model
y_train_final = y_encoder.fit_transform(y_train.values.reshape(-1, 1)).ravel()

# Display the encoded training outcome variable
y_train_final
```

```
[98]: array([1., 1., 1., ..., 1., 1., 0.])
```

3.0.7 Task 3d. Model building

Construct a model and fit it to the training set.

```
[99]: # Construct a logistic regression model and fit it to the training set
      clf = LogisticRegression(random_state=0,max_iter=800).
      ↪fit(X_train_final,y_train_final)
```

3.1 PACE: Execute

Consider the questions in your PACE Strategy Document to reflect on the Execute stage.

3.1.1 Taks 4a. Results and evaluation

Evaluate your model.

Encode categorical features in the testing set using an appropriate method.

```
[100]: # Select the testing features that needs to be encoded
      X_test_to_encode = X_test[["claim_status", "author_ban_status"]]

      # Display first few rows
      X_test_to_encode.head()
```

```
[100]:      claim_status  author_ban_status
      21061      opinion      active
      31748      opinion      active
      20197      claim      active
      5727      claim      active
      11607      opinion      active
```

```
[101]: # Transform the testing features using the encoder
      X_test_encoded = X_encoder.transform(X_test_to_encode)

      # Display first few rows of encoded testing features
      X_test_encoded
```

```
[101]: array([[1., 0., 0.],
      [1., 0., 0.],
      [0., 0., 0.],
      ...,
      [1., 0., 0.],
      [0., 0., 1.],
      [1., 0., 0.]])
```

```
[102]: # Place encoded testing features (which is currently an array) into a dataframe
      X_test_encoded_df = pd.DataFrame(data=X_test_encoded, columns=X_encoder.
      ↪get_feature_names_out())
```

```
# Display first few rows
X_test_encoded_df.head()
```

```
[102]:      claim_status_opinion  author_ban_status_banned  \
0                1.0                0.0
1                1.0                0.0
2                0.0                0.0
3                0.0                0.0
4                1.0                0.0

      author_ban_status_under review
0                0.0
1                0.0
2                0.0
3                0.0
4                0.0
```

```
[103]: # Display first few rows of `X_test` with `claim_status` and
      ↪ `author_ban_status` columns dropped (since these features are being
      ↪ transformed to numeric)
      ### YOUR CODE HERE ###
      X_test.drop(columns=["claim_status", "author_ban_status"]).head()
```

```
[103]:      video_duration_sec  video_view_count  video_share_count  \
21061                41            2118.0            57.0
31748                27            5701.0            157.0
20197                31           449767.0          75385.0
5727                 19           792813.0          56597.0
11607                54            2044.0             68.0

      video_download_count  video_comment_count
21061                5.0                2.0
31748                1.0                0.0
20197           5956.0            728.5
5727           5146.0            728.5
11607           19.0                2.0
```

```
[104]: # Concatenate `X_test` and `X_test_encoded_df` to form the final dataframe for
      ↪ training data (`X_test_final`)
      X_test_final = pd.concat([X_test.drop(columns=["claim_status",
      ↪ "author_ban_status"]), reset_index(drop=True), X_test_encoded_df], axis=1)

      # Display first few rows
      X_test_final.head()
```

```
[104]:      video_duration_sec  video_view_count  video_share_count  \
0                41            2118.0            57.0
```

1	27	5701.0	157.0
2	31	449767.0	75385.0
3	19	792813.0	56597.0
4	54	2044.0	68.0

	video_download_count	video_comment_count	claim_status_opinion \
0	5.0	2.0	1.0
1	1.0	0.0	1.0
2	5956.0	728.5	0.0
3	5146.0	728.5	0.0
4	19.0	2.0	1.0

	author_ban_status_banned	author_ban_status_under review
0	0.0	0.0
1	0.0	0.0
2	0.0	0.0
3	0.0	0.0
4	0.0	0.0

Test the logistic regression model. Use the model to make predictions on the encoded testing set.

```
[105]: # Use the logistic regression model to get predictions on the encoded testing set
      ↪set
      y_pred = clf.predict(X_test_final)
```

Display the predictions on the encoded testing set.

```
[106]: # Display the predictions on the encoded testing set
      y_pred
```

```
[106]: array([1., 1., 0., ..., 1., 0., 1.])
```

Display the true labels of the testing set.

```
[107]: # Display the true labels of the testing set
      y_test
```

```
[107]: 21061      verified
      31748      verified
      20197      verified
      5727  not verified
      11607  not verified
      ...
      14756  not verified
      26564      verified
      14800  not verified
      35705      verified
      31060      verified
```

Name: verified_status, Length: 8942, dtype: object

Encode the true labels of the testing set so it can be compared to the predictions.

```
[108]: # Encode the testing outcome variable
#       - Adjusting the shape of `y_test` before passing into `.transform()`, since
#         ↳ it takes in 2D array
#       - Using `.ravel()` to flatten the array returned by `.transform()`, so that
#         ↳ it can be used later to compare with predictions
y_test_final = y_encoder.transform(y_test.values.reshape(-1, 1)).ravel()

# Display the encoded testing outcome variable
y_test_final
```

```
[108]: array([1., 1., 1., ..., 0., 1., 1.])
```

Confirm again that the dimensions of the training and testing sets are in alignment since additional features were added.

```
[109]: # Get shape of each training and testing set
X_train_final.shape, y_train_final.shape, X_test_final.shape, y_test_final.shape
```

```
[109]: ((26826, 8), (26826,), (8942, 8), (8942,))
```

3.1.2 Task 4b. Visualize model results

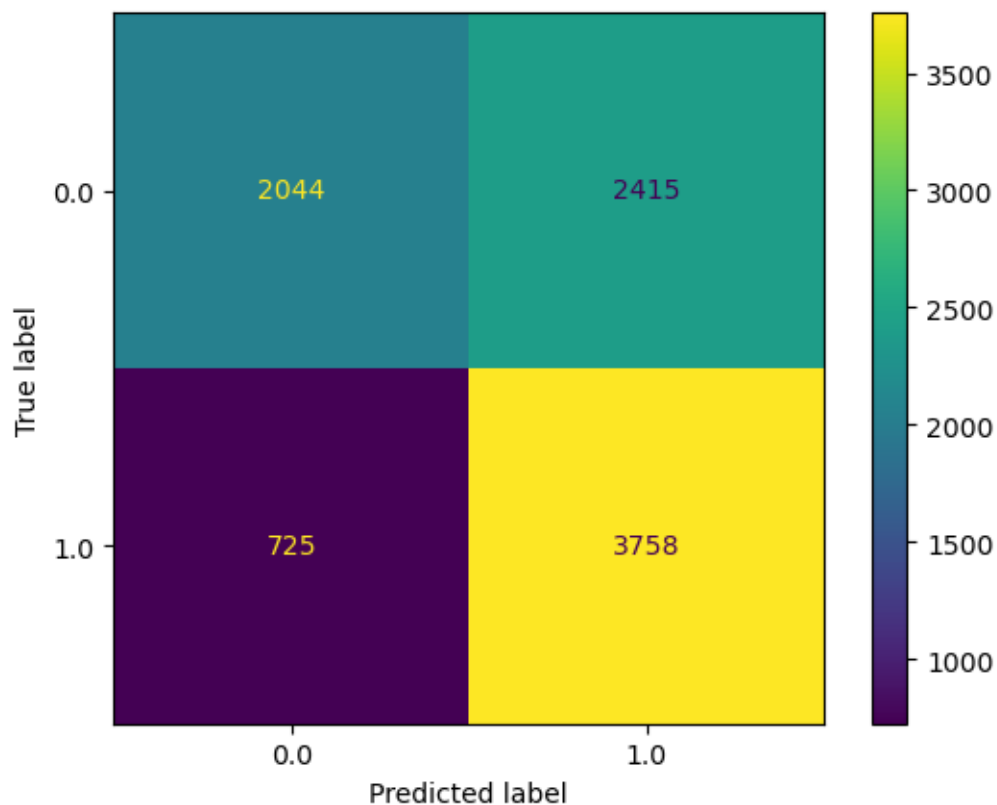
Create a confusion matrix to visualize the results of the logistic regression model.

```
[111]: # Compute values for confusion matrix
log_cm = confusion_matrix(y_test_final, y_pred, labels=clf.classes_)

# Create display of confusion matrix
log_disp = ConfusionMatrixDisplay(confusion_matrix=log_cm, display_labels=clf.
↳ classes_)

# Plot confusion matrix
log_disp.plot()

# Display plot
plt.show()
```

Create a classification report that includes precision, recall, f1-score, and accuracy metrics to evaluate the performance of the logistic regression model.

```
[113]: # Create a classification report
target_labels = ["verified", "not verified"]
print(classification_report(y_test_final, y_pred, target_names=target_labels))
```

	precision	recall	f1-score	support
verified	0.74	0.46	0.57	4459
not verified	0.61	0.84	0.71	4483
accuracy			0.65	8942
macro avg	0.67	0.65	0.64	8942
weighted avg	0.67	0.65	0.64	8942

3.1.3 Task 4c. Interpret model coefficients

```
[115]: # Get the feature names from the model and the model coefficients (which
      ↪ represent log-odds ratios)
      # Place into a DataFrame for readability
      pd.DataFrame(data={"Feature Name":clf.feature_names_in_, "Model Coefficient":
      ↪clf.coef_[0]})
```

```
[115]:
```

	Feature Name	Model Coefficient
0	video_duration_sec	8.607893e-03
1	video_view_count	-2.132079e-06
2	video_share_count	5.930971e-06
3	video_download_count	-1.099775e-05
4	video_comment_count	-6.404235e-04
5	claim_status_opinion	3.908384e-04
6	author_ban_status_banned	-1.781741e-05
7	author_ban_status_under review	-9.682447e-07

3.1.4 Task 4d. Conclusion

Key takeaways:

- The dataset has a few strongly correlated variables, which might lead to multicollinearity issues when fitting a logistic regression model. We decided to drop `video_like_count` from the model building.
- Based on the logistic regression model, each additional second of the video is associated with 0.009 increase in the log-odds of the user having a verified status.
- The logistic regression model had not great, but acceptable predictive power: a precision of 61% is less than ideal, but a recall of 84% is very good. Overall accuracy is towards the lower end of what would typically be considered acceptable.

We developed a logistic regression model for verified status based on video features. The model had decent predictive power. Based on the estimated model coefficients from the logistic regression, longer videos tend to be associated with higher odds of the user being verified. Other video features have small estimated coefficients in the model, so their association with verified status seems to be small.